

# git

it

# objectives

By the end of this workshop, you will understand (some of)

**what git does for you  
behind the scenes**

# objectives

By the end of this workshop, you will be able to

**build commits that represent  
meaningful changes to the  
codebase**

# objectives

By the end of this workshop, you will be able to

**make the right decisions when  
rebasing or merging**

# objectives

By the end of this workshop, you will be able to

**efficiently find the cause of  
regressions with** `bisect`

# not covered

- working with remote repositories


# not covered

- git for mercurial / subversion / cvs users

# not covered

- anything github-specific

(sorry)





# not covered

- *basics* like creating branches, cloning, pushing, pulling...

# topics

- the most likely reason we use git
- merging & rebasing
- conflict resolution w/ diff3
- finding the cause of bugs with bisect
- cool git tips

# topics

- the most likely reason we use git
- merging & rebasing
- conflict resolution w/ diff3
- finding the cause of bugs with bisect
- cool git tips

# why has git taken over open source?

- branches are first-class objects
- distributed nature lends itself well to forking
- free
- really really fast

# why is git so fast?

- ~~branches are first-class objects~~
- ~~distributed nature lends itself well to forking~~
- ~~free~~
- really really fast – *people use fast things*  
(often at the sacrifice of usability)

# why is git so fast?

git stores and manages the data that you give it in a unique and efficient way

# why is git so fast?

But first, some background

The old-school way:

CVS, subversion, mercurial, etc.,  
all fell down this trap:

Differences between revisions are  
therefore...

Only store the changes between revisions

source dungeon, etc,

are very small



*“What a great idea!”*  
(Dick Gruene, creator of CVS)

# storing diffs per-re

Storing diffs between revisions has *got* to be optimal ...after all, diffs are usually pretty small compared to the entire codebase.



SHUT THE  
F@#\$ UP!  
You've got it all  
wrong! Crawl back  
into your hole in the  
ground you  
**caveman!**





# storing diffs per-revision

While diffing between revisions is efficient on disk, it is **not efficient** for doing most of the important things that a good version control system should do.

Switching between branches takes a big performance hit because of this.

# an example

What follows is an example of why storing diffs per revision storage scheme is inefficient when switching between branches.

# not so fast... switching

*When you do this:*

```
svn switch myBranch
```

*svn, cvs, hg, etc... do this:*

finds the common ancestor of `myBranch` and your current branch, walks down to from your current branch, then walks to `myBranch`, recording the changes along the way, then applies those changes to your current working directory. –  $O(n)$

*When you do this:*

```
git checkout myBranch
```

*git does this:*


finds files **that already exist** in `.` `git/objects` that are in `myBranch` and uses them to replace files in your current working directory. –  $O(1)$

*way faster! awesome!*

# fast switching!


*When you do this:*

**svn switch myBranch**



*When you do this:*

**git checkout myBranch**



*in other words:*

git doesn't need to care  
about where your revision  
came from in order to  
switch to it.

# how does git pull this off?

git does not need to walk the history graph of your repository to apply changes.

git organizes the files in your repository so it will have near-immediate access to every revision of every file.

# how does git pull this off?

Files are keyed by their SHA-1 hash in `.git/objects`

Commits as stored by git (indirectly) store the path and SHA-1 hash of each file in your repository.

# what about space?

That is very wasteful of  
disk space; I'd never  
design a version control  
system that worked like  
that!



# what about space?

It turns out that with some clever pattern matching, all of your repository's data can be stored as efficiently *without* basing the diffs on the history graph.



# what about space?

git uses a structure called a “packfile” to efficiently store similar files.

packfiles are stored in `.git/objects` alongside all the other files in your repository, but you don't (usually) interact with them directly.

# packfiles!

packfiles are usually created when git does its garbage collection.

You've probably noticed garbage collection happening occasionally when you `git pull` or `git fetch` (or run `git gc` manually.)

# packfiles!

packfile creation and optimization is one of the things that is happening when you see this:

```
thomashallocks-MacBook-Pro~DC(develop|...) % g
```

# packfiles!

Have a look at the “internals” chapter of the git book. It explains in very clear language how git keeps track of absolutely everything:

<https://git-scm.com/book/en/Git-Internals>

# ok, git is fast. now what?

git's speed opens up a slew of great features:

- rebasing
- branches as first-class objects
- `commit --amend`
- its distributed nature
- quickly comparing files across arbitrary revisions

# ok, git is fast. now what?

git's speed opens up a slew of great features:

- **rebasing**
- branches as first-class objects
- `commit --amend`
- its distributed nature
- quickly comparing files across arbitrary revisions

# rebasing

At the very least, you're probably familiar with the term "rebasing."

Most of you have been using it in some capacity:

- `git pull --rebase`
- `git rebase origin/master`

# rebasing

...but what exactly is “rebasing”?


re – change

base – a structure on which something draws or depends

When you rebase, you are changing the parent commit of your branch. In order to accommodate this change, all the commits after the parent commit must be reconstructed.




# rebasing (and merging)



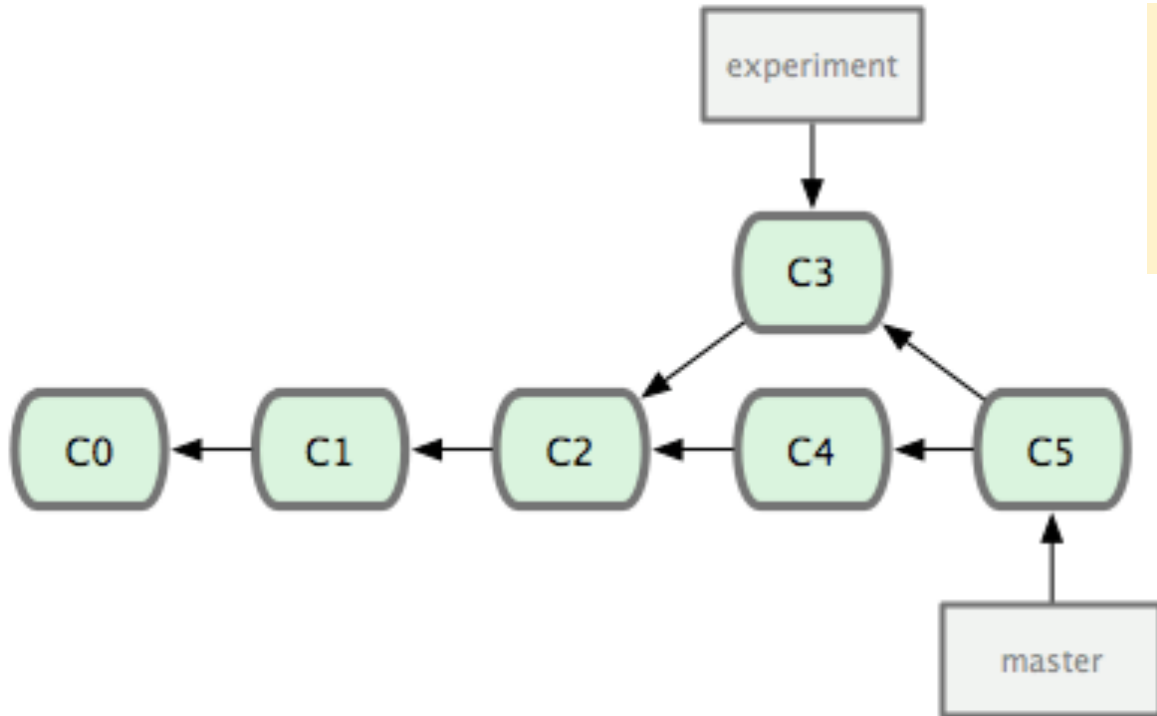
**Consider this situation**

# rebasing (and merging)




The easiest way to integrate `experiment` into `master` is to use the `merge` command. Merge performs a three-way merge between the two latest branch snapshots (C3 and C4) and the most recent common ancestor of the two (C2), creating a new snapshot (and commit)

# rebasing (and merging)



C5 is the  
merge commit

# rebasing (and merging)




However, there is another way:

You can take the *patch* of the change that was introduced in C3 and re-apply it *on top of* C4. In Git, this is called *rebasing*.

With the `rebase` command, git will take all unique changes on one branch and replay them on another one.

# rebasing (and merging)




```
$ git checkout experiment
```

```
$ git rebase master
```

First, rewinding head to  
replay your work on top  
of it...


Applying: added staged  
command

# rebasing (and merging)



There is no difference between the code in the rebased version and the code in the merged version in the prior slide.

# rebasing (and merging)



However, the history is cleaner in the rebased version because it does not contain a merge commit.

# changing the parent branch

Rebasing is great for producing clean history, but it can also take your branch places that a merge can not.


Change the parent of your branch like this:

```
git rebase --onto <new parent> <old parent>
```




# moving to the release branch

Here's a relevant scenario for us at 1stdibs. Consider that you've been keeping your `feature` branch up-to-date with `master` on bunsen:



# moving to the release branch


You let it slip that you've finished `feature` and product decides to release it early, so now just the commits unique to `feature` have to be included in `v5.33`. In the graph below, that would be D and E, but *not* B or C.



point at which  
master was cut


# moving to the release branch

You can't just merge `feature` as it stands into `v5.33` because doing so would also include commits B and C, which aren't scheduled for `v5.33`.



# moving to the release branch


So, you just want to get D and E into v5.33. In a simple scenario like this, you could get away with cherry-picking D and E into a new branch of v5.33, but if `feature` had 30 unique commits, you'd have a lot of repetitive work to do.



# moving to the release branch

```
git rebase --onto v5.33
```


master



# moving to the release branch


```
git rebase --onto v5.33
```

master



# don't rebase shared branches

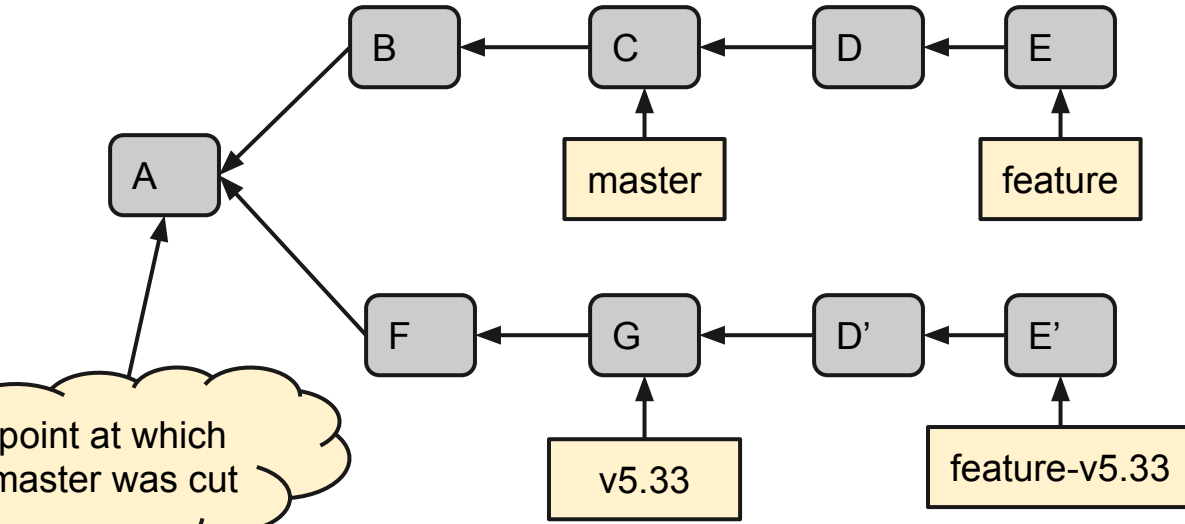
If you're sharing this branch with someone else, do not push your rebased branch to your remote under the same name.



# don't rebase shared branches

Instead, rename your rebased branch before you push it:


```
git branch -m feature-v5.33
```





# moving to the release branch

When your coworker is good and ready, they can move their commits from `feature` to `feature-v5.33` in the same way that you moved `feature`'s commits from `master` onto `v5.33` with `rebase --onto`



# the perils of rebasing

The git book presents a common scenario that arises when multiple people work off rebased branches:

<http://git-scm.com/book/en/v2/Git-Branching-Rebasing#The-Perils-of-Rebasing>

# git rebase -i <parent>

run `rebase` with the `-i` flag and git will show you in your editor the commits that it plans to rebase.

You can edit this plan...

# git rebase -i <parent>

node\_modules/bunsen/.git/rebase-merge / git-rebase-todo

```
1 pick f0af820 MCPERSON-1178 MCPERSON-1188: add private offer indicator to semi
2 pick d1b7627 MCPERSON-1223 Hide private offer faq
3 pick f5e3636 GOODS-1102 Revert image path validation
4 pick 702849c added tooltip and new version of icons
5 pick 4e5127f icons added
6 pick 54dfb75 MCPERSON-1218 fixing width of offer column
7 pick 05b8907 MCPERSON-1227 Swap font awesome offer icon with bunsen icon
8 pick e8c6d89 MCPERSON-1182 excluding adverbs from inline dates
9 pick 5bcbcb8 MCPERSON-1183 when checking for canceled transactions, take into
10 pick 7d5bee8 MCPERSON-1227 Private offer amount inbox spacing
11 pick 4ebafe3 MONEY-1650 can buy no price items with private offer
12 pick 281a38b SHIPALYTIC-1027 remove endpoint using queryservice for vertical
13 pick 0a947e9 MONEY-1651 - change computeAdjustments back to setEachModel
14 pick 5e9c49e MCPERSON-1195 looking at transaction's status to determine if pr
15 pick 3e8900e MONEY-1245 don't add fetchOptions to transaction model
16 pick 1699824 removed lint
17 pick 4c9f14f GOODS-1677 add some methods to Publication model
18 pick fcc4e77 GOODS-1692 - updating publication models
19 pick 3762c5a GOODS-1851 replaced publishOptions calls
20 pick 4b4a4ee Bringing back setPublishOption
21
22 # Rebase baa84bf..4b4a4ee onto baa84bf
23 #
24 # Commands:
25 # p, pick = use commit
26 # r, reword = use commit, but edit the commit message
27 # e, edit = use commit, but stop for amending
28 # s, squash = use commit, but meld into previous commit
29 # f, fixup = like "squash", but discard this commit's log message
```

You can remove, reorder or even add more commits to this list.

You can also meld a commit with the one above it by changing the `pick` to `fixup`.

If you want to keep the commit message of both messages, use `squash` instead.

```
git rebase -i <parent>
```

*...live demo in  
progress...*

**recommended  
git configs**





Have a look at my .gitconfig:

[github.com/antialias/dotfiles/blob/master/.gitconfig](https://github.com/antialias/dotfiles/blob/master/.gitconfig)

# .gitconfig

```
[rebase]
  autosquash = true
```

--autosquash by default when doing  
interactive rebase



# .gitconfig

```
[rerere]  
    enabled = true
```

or just run this command right now (I'll wait)

```
git config --global rerere.enabled true
```

# .gitconfig

```
[rerere]
```

```
enabled = true
```

***reuse recorded resolution*** of conflicted merges

Git will remember how you resolved every conflict, and will apply that same resolution if it sees the same conflict in the future.

# .gitconfig

```
[alias]
    lg = log --color --graph --pretty=format:'%Cred%h%
Creset -%C(yellow)%d%Creset %s %Cgreen(%cr) %C(bold blue)<%
an>%Creset' --abbrev-commit
```

aliases git lg to show a really nicely formatted log instead of the overly-verbose one that git shows by default.

# .gitconfig

```
thomashallocks-MacBook-Pro-B(demo-fixup|📌) % git lg -30
* 4b4a4ee - (HEAD, tom/demo-fixup, demo-fixup) Bringing back setPublishOption (38 minutes ago)<prestifidibs>
* 3762c5a - GOODS-1851 replaced publishOptions calls (38 minutes ago)<Rich>
* fcc4e77 - GOODS-1692 - updating publication models (38 minutes ago)<prestifidibs>
* 4c9f14f - GOODS-1677 add some methods to Publication model (38 minutes ago)<donfusilli>
* 1699824 - removed lint (38 minutes ago)<Thomas Hallock>
* 3e8900e - MONEY-1245 don't add fetchOptions to transaction model (38 minutes ago)<Augusto Corvalan>
* 5e9c49e - MCPERSON-1195 looking at transaction's status to determine if private offer transaction show be consider
* 0a947e9 - MONEY-1651 - change computeAdjustments back to setEachModel (4 hours ago)<Rob Richard>
* 281a38b - SHIPALYTIC-1027 remove endpoint using queryservice for vertical categories (4 hours ago)<Joseph Ahn>
* 4ebafe3 - MONEY-1650 can buy no price items with private offer (4 hours ago)<Rob Richard>
* 7d5bee8 - MCPERSON-1227 Private offer amount inbox spacing (4 hours ago)<Peter Jodogne>
* 5bcbcbb - MCPERSON-1183 when checking for canceled transactions, take into account case where item was marked sold
r items sold through checkout (4 hours ago)<Brett Ashford>
* e8c6d89 - MCPERSON-1182 excluding adverbs from inline dates (4 hours ago)<Thomas Hallock>
* 05b8907 - MCPERSON-1227 Swap font awesome offer icon with bunsen icon (4 hours ago)<Peter Jodogne>
* 54dfb75 - MCPERSON-1218 fixing width of offer column (4 hours ago)<Thomas Hallock>
* 4e5127f - icons added (4 hours ago)<Liam Dickson>
* 702849c - (refs/bisect/good-702849ca1563a2d38683e6df6c6b88d0083de3fc) added tooltip and new version of icons (4 ho
* f5e3636 - GOODS-1102 Revert image path validation (4 hours ago)<takeWarning>
* d1b7627 - MCPERSON-1223 Hide private offer faq (4 hours ago)<Peter Jodogne>
* f0af820 - MCPERSON-1178 MCPERSON-1188: add private offer indicator to semi-collapsed message view and apply inacti
* baa84bf - MONEY-1633 buyer offer takes a financial amount (4 hours ago)<Rob Richard>
* 5183f40 - MCPERSON-1155 move normalizeCurrencyFormat() into bunsen/numberConversions.js so it can be used in admin
* f0f8f73 - MCPERSON-1220 showing actual status for private offer headsip when we don't have any predefined logic. (
* d0afbfo - GOODS-1572 upcoming arrow should have white fill, not default black fill (4 hours ago)<donfusilli>
* b3356f3 - MONEY-1245 added disabled-action color var (4 hours ago)<Augusto Corvalan>
```

# .gitconfig

```
thomashallocks-MacBook-Pro-B(master|✓) % git lg -30
* 954e251 - (HEAD, origin/master, origin/HEAD, master) Merge pull request #1112 from jakemccloskey/mixin-helper (31 minutes ago)<dale tan>
| \
| * 6aebd27 - Added mixin helper for ES2015 classes (8 days ago)<Jake McCloskey>
| * | 6412847 - Bringing back setPublishOption (24 hours ago)<prestifidibs>
| * | f4f317e - Merge pull request #1110 from RichLandy/feature-publishing-statuses (25 hours ago)<prestifidibs>
| \ \
| * | 10139fa - GOODS-1851 replaced publishOptions calls (25 hours ago)<Rich>
| * | | f5d73bf - Merge pull request #1141 from prestifidibs/feature-listings (25 hours ago)<prestifidibs>
| \ \ \
| \ / \
| / / \
| * | f34fc62 - GOODS-1692 - updating publication models (26 hours ago)<prestifidibs>
| / /
| * | 68d0b9f - Merge pull request #1127 from donfusilli/feature-publication-model-additions (28 hours ago)<prestifidibs>
| \ \
| * | 1cafdd3 - GOODS-1677 add some methods to Publication model (28 hours ago)<donfusilli>
| / /
| * | b66cba3 - removed lint (4 days ago)<Thomas Hallock>
| * | 6a2226d - Merge branch 'v5.30' of https://github.com/1stdibs/bunsen (4 days ago)<Rob Richard>
| \ \
| * \ 7fef954 - (origin/v5.30, origin/integration-private-offer-status, bre/integration-private-offer-status, v5.30, integration-private-offer-status-awaiting-manual-review (5 days ago)<brettjashford>
| | \ \
| | * | 15dc972 - (tom/bugfix-consider-status-awaiting-manual-review, bugfix-consider-status-awaiting-manual-review) MCPERSON-1195 looking at the show be considered 'accepted' (5 days ago)<Thomas Hallock>
| | * | bc387a1 - fixup! MCPERSON-1220 showing actual status for private offer heads up when we don't have any predefined logic. (5 days ago)<Thomas Hallock>
| | / /
| * | 8234fbd - Merge pull request #1126 from augustocorvalan/feature-saved-address (5 days ago)<Rob Richard>
| \ \ \
| * | | ab307cb - MONEY-1245 don't add fetchOptions to transaction model (5 days ago)<Augusto Corvalan>
```

# .gitconfig

```
[merge]  
    conflictstyle = diff3
```

shows the common ancestor between the “ours” and “theirs” blocks of a conflict.

# .gitconfig

cauliflower

<<<<<< HEAD

peas

potatoes

||||||| merged common ancestors

peas

=====

>>>>>> topic

tomatoes

# .gitconfig

cauliflower

<<<<<< HEAD

peas

potatoes

=====

>>>>>> topic

tomatoes



# .gitconfig

```
[merge]  
    conflictstyle = diff3
```

PHPStorm doesn't show the common ancestor block, so you'll have to resolve the conflict directly in your source file. Learn to read it.

# .gitconfig

```
[merge]  
    conflictstyle = diff3
```

More information here:

<http://psung.blogspot.com/2011/02/reducing-merge-headaches-git-meets.html>

**end of slides - *demo time!***

# bisect

*John Rodriguez, Sr. Front-End Engineer*

will now demonstrate how to find the cause of regressions by using a clever git tool called

bisect

**caution: extra slides ahead**


boredom surely awaits. you have been warned

# packfiles!

git's algorithm for creating a packfile goes something like this:


1. find similar files in `.git/objects`
2. store the differences between them in a special file
3. remove all but the base file from which the differences are computed.
4. list the files and the diff files in the packfile, with the order of the files indicating how the files relate to each other.

# changing the parent branch



Suppose that the code for `client` is ready to be integrated with `master`, but the code for `server` is not.


# changing the parent branch



`rebase --onto` can be used to integrate the code for `client` into `master` while excluding the code for `server`.




# changing the parent branch



checkout client


rebase --onto master server

# time for a fast-forward




Now we can fast-forward  
master up to client with  
checkout master  
merge client

# time for a fast-forward



Now we can fast-forward  
master up to client with  
checkout master  
merge client

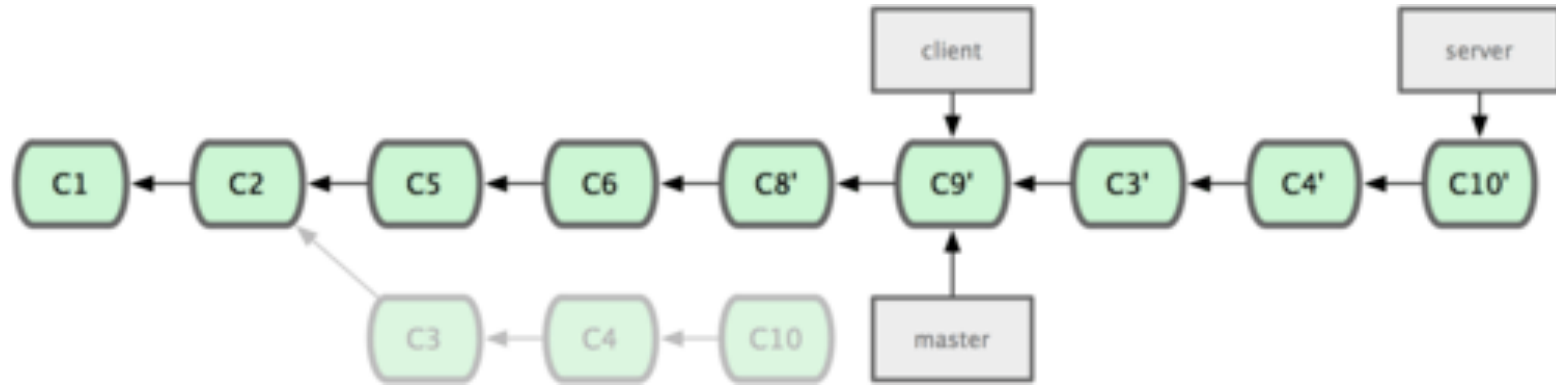
# going all the way



Suppose `server` is now ready to be integrated into `master`. We can rebase it on top of `master` like this:

```
checkout server  
rebase master
```


# completely linear history



Now, we can merge server into master with:

```
git checkout master  
git merge server
```

# completely linear history



Again, git figures out that no merge commit is needed so it does a fast-forward merge of `server` into `master`.